

# Full-Stack RAG with Local LLM

Laboratory Manual

CSI403 - Full Stack Development

Faculty of Information Technology, Sripatum University

Academic Year 2568

# Contents

<b>1 Lab 1: Git + Python Fundamentals (3.75%)</b>	<b>3</b>
1.1 Objectives . . . . .	3
1.2 Tasks . . . . .	3
1.2.1 Task 1: Fork Repository . . . . .	3
1.2.2 Task 2: Clone to Local . . . . .	3
1.2.3 Task 3: Setup Environment . . . . .	3
1.2.4 Task 4: Create Feature Branch . . . . .	3
1.2.5 Task 5: Write Python OOP Example . . . . .	3
1.2.6 Task 6: Commit and Push . . . . .	4
1.2.7 Task 7: Create Pull Request . . . . .	4
1.3 Deliverables . . . . .	4
<b>2 Lab 2: Docker + OpenSearch (3.75%)</b>	<b>5</b>
2.1 Objectives . . . . .	5
2.2 Tasks . . . . .	5
2.2.1 Task 1: Install Docker Desktop . . . . .	5
2.2.2 Task 2: Run OpenSearch . . . . .	5
2.2.3 Task 3: Verify . . . . .	5
2.2.4 Task 4: Setup Hybrid Search Pipeline . . . . .	5
2.2.5 Task 5: Check Health . . . . .	5
2.3 Deliverables . . . . .	6
<b>3 Lab 3: FastAPI (3.75%)</b>	<b>7</b>
3.1 Objectives . . . . .	7
3.2 Tasks . . . . .	7
3.2.1 Task 1: Study api.py . . . . .	7
3.2.2 Task 2: Run Server . . . . .	7
3.2.3 Task 3: Test Endpoints . . . . .	7
3.2.4 Task 4: Access Swagger Docs . . . . .	7
3.2.5 Task 5: Add New Endpoint . . . . .	7
3.3 Repository . . . . .	7
3.4 Deliverables . . . . .	8
<b>4 Lab 4: OpenSearch Integration (3.75%)</b>	<b>9</b>
4.1 Objectives . . . . .	9
4.2 Tasks . . . . .	9
4.2.1 Task 1: Connect to OpenSearch . . . . .	9
4.2.2 Task 2: Create Index . . . . .	9
4.2.3 Task 3: Index Documents . . . . .	9
4.2.4 Task 4: Implement Search . . . . .	9
4.3 Repository . . . . .	10
4.4 Deliverables . . . . .	10
<b>5 Lab 5: Embeddings (3.75%)</b>	<b>11</b>
5.1 Objectives . . . . .	11
5.2 Tasks . . . . .	11
5.2.1 Task 1: Study embedding.py . . . . .	11

5.2.2	Task 2: Run Indexing . . . . .	11
5.2.3	Task 3: Verify . . . . .	11
5.2.4	Task 4: Add New Documents . . . . .	11
5.2.5	Task 5: Re-index . . . . .	11
5.2.6	Task 6: Test Search . . . . .	11
5.3	Repository . . . . .	11
5.4	Deliverables . . . . .	12
<b>6</b>	<b>Lab 6: RAG + Local LLM + Streamlit (3.75%)</b>	<b>13</b>
6.1	Objectives . . . . .	13
6.2	Tasks . . . . .	13
6.2.1	Task 1: Install Ollama . . . . .	13
6.2.2	Task 2: Pull Model . . . . .	13
6.2.3	Task 3: Test LLM . . . . .	13
6.2.4	Task 4: Run API with LLM . . . . .	13
6.2.5	Task 5: Run Streamlit . . . . .	13
6.2.6	Task 6: Test Complete Flow . . . . .	13
6.2.7	Task 7: Modify Prompt . . . . .	13
6.3	Repository . . . . .	13
6.4	Deliverables . . . . .	14
<b>7</b>	<b>Lab 7: Docker Compose (3.75%)</b>	<b>15</b>
7.1	Objectives . . . . .	15
7.2	Tasks . . . . .	15
7.2.1	Task 1: Clone Repository . . . . .	15
7.2.2	Task 2: Study docker-compose.yml . . . . .	15
7.2.3	Task 3: Start Services . . . . .	15
7.2.4	Task 4: Check Status . . . . .	15
7.2.5	Task 5: View Logs . . . . .	15
7.2.6	Task 6: Test Application . . . . .	15
7.2.7	Task 7: Stop Services . . . . .	15
7.3	Repository . . . . .	15
7.4	Deliverables . . . . .	16
<b>8</b>	<b>Lab 8: CI/CD + Testing (3.75%)</b>	<b>17</b>
8.1	Objectives . . . . .	17
8.2	Tasks . . . . .	17
8.2.1	Task 1: Write test_api.py . . . . .	17
8.2.2	Task 2: Run Tests . . . . .	17
8.2.3	Task 3: Create Jenkinsfile . . . . .	17
8.2.4	Task 4: Create GitHub Actions . . . . .	18
8.3	Repository . . . . .	18
8.4	Deliverables . . . . .	18
<b>Summary: All Labs</b>		<b>19</b>

# 1 Lab 1: Git + Python Fundamentals (3.75%)

## 1.1 Objectives

- Fork and clone Generic-RAG repository
- Practice Git workflow
- Review Python fundamentals

## 1.2 Tasks

### 1.2.1 Task 1: Fork Repository

1. Go to <https://github.com/amornpan/Generic-RAG>
2. Click “Fork” button
3. You now have your own copy

### 1.2.2 Task 2: Clone to Local

```
git clone https://github.com/YOUR_USERNAME/Generic-RAG.git  
cd Generic-RAG
```

### 1.2.3 Task 3: Setup Environment

```
conda create -n rag_env python=3.10 -y  
conda activate rag_env  
pip install -r requirements.txt
```

### 1.2.4 Task 4: Create Feature Branch

```
git checkout -b feature/lab01-python
```

### 1.2.5 Task 5: Write Python OOP Example

Create lab01\_example.py:

```
class Document:  
    def __init__(self, title: str, content: str):  
        self.title = title  
        self.content = content  
  
    def get_summary(self) -> str:  
        return self.content[:100] + "..."  
  
# Test  
doc = Document("Test", "This is a test document...")  
print(doc.get_summary())
```

### 1.2.6 Task 6: Commit and Push

```
git add .
git commit -m "Lab 1: Add Python OOP example"
git push origin feature/lab01-python
```

### 1.2.7 Task 7: Create Pull Request

1. Go to GitHub
2. Click “Compare & pull request”
3. Create PR

## 1.3 Deliverables

Forked repository

Feature branch created

Python file committed

Pull Request created

Screenshot of PR

**Deadline:** Sunday 23:59

## 2 Lab 2: Docker + OpenSearch (3.75%)

### 2.1 Objectives

- Install Docker Desktop
- Run OpenSearch container
- Configure Hybrid Search Pipeline

### 2.2 Tasks

#### 2.2.1 Task 1: Install Docker Desktop

1. Download from <https://docker.com>
2. Install and start
3. Verify: `docker -version`

#### 2.2.2 Task 2: Run OpenSearch

```
docker run -d --name opensearch-node \
-p 9200:9200 -p 9600:9600 \
-e "discovery.type=single-node" \
-e "DISABLE_SECURITY_PLUGIN=true" \
opensearchproject/opensearch:2.11.1
```

#### 2.2.3 Task 3: Verify

```
curl http://localhost:9200
```

#### 2.2.4 Task 4: Setup Hybrid Search Pipeline

```
curl -X PUT "localhost:9200/_search/pipeline/hybrid-search-pipeline" \
-H "Content-Type: application/json" \
-d '{
  "phase_results_processors": [
    {
      "normalization-processor": {
        "normalization": {"technique": "min_max"},
        "combination": {
          "technique": "arithmetic_mean",
          "parameters": {"weights": [0.3, 0.7]}
        }
      }
    }
  ],
  "script"
}'
```

#### 2.2.5 Task 5: Check Health

```
curl http://localhost:9200/_cluster/health?pretty
```

## 2.3 Deliverables

Docker installed

OpenSearch running

Pipeline created

Screenshots

**Deadline:** Sunday 23:59

## 3 Lab 3: FastAPI (3.75%)

### 3.1 Objectives

- Study api.py from Generic-RAG
- Run FastAPI server
- Test API endpoints

### 3.2 Tasks

#### 3.2.1 Task 1: Study api.py

Review the structure:

- FastAPI app
- Pydantic models
- Endpoints: /health, /search, /query

#### 3.2.2 Task 2: Run Server

```
cd Generic-RAG
conda activate rag_env
python api.py
```

#### 3.2.3 Task 3: Test Endpoints

```
# Health check
curl http://localhost:9000/health

# Search
curl -X POST http://localhost:9000/search \
-H "Content-Type: application/json" \
-d '{"query": "test", "top_k": 5}'
```

#### 3.2.4 Task 4: Access Swagger Docs

Open: <http://localhost:9000/docs>

#### 3.2.5 Task 5: Add New Endpoint

Add /documents endpoint to list all documents.

### 3.3 Repository

<https://github.com/amornpan/Generic-RAG>

### 3.4 Deliverables

Server running

Endpoints tested

New endpoint added

Screenshots

**Deadline:** Sunday 23:59

## 4 Lab 4: OpenSearch Integration (3.75%)

### 4.1 Objectives

- Connect to OpenSearch from Python
- Implement hybrid search
- Index and search documents

### 4.2 Tasks

#### 4.2.1 Task 1: Connect to OpenSearch

```
from opensearchpy import OpenSearch

client = OpenSearch(
    hosts=[{"host": "localhost", "port": 9200}],
    use_ssl=False
)
print(client.info())
```

#### 4.2.2 Task 2: Create Index

```
index_body = {
    "settings": {"index": {"knn": True}},
    "mappings": {
        "properties": {
            "content": {"type": "text"},
            "content_vector": {
                "type": "knn_vector",
                "dimension": 1024
            }
        }
    }
}
client.indices.create(index="documents", body=index_body)
```

#### 4.2.3 Task 3: Index Documents

```
doc = {
    "content": "Sample document text",
    "content_vector": [0.1, 0.2, ...] # 1024 dims
}
client.index(index="documents", body=doc)
```

#### 4.2.4 Task 4: Implement Search

Implement hybrid search combining vector and BM25.

### **4.3 Repository**

<https://github.com/amornpan/Generic-RAG>

### **4.4 Deliverables**

Connection working

Index created

Documents indexed

Search implemented

**Deadline:** Sunday 23:59

## 5 Lab 5: Embeddings (3.75%)

### 5.1 Objectives

- Run embedding.py
- Understand chunking
- Add new documents

### 5.2 Tasks

#### 5.2.1 Task 1: Study embedding.py

Review the pipeline:

1. Load documents from md\_corpus/
2. Chunk documents
3. Create embeddings (bge-m3)
4. Index to OpenSearch

#### 5.2.2 Task 2: Run Indexing

```
cd Generic-RAG  
conda activate rag_env  
python embedding.py
```

#### 5.2.3 Task 3: Verify

```
curl http://localhost:9200/documents/_count
```

#### 5.2.4 Task 4: Add New Documents

Add 3 new .md files to md\_corpus/ folder.

#### 5.2.5 Task 5: Re-index

```
python embedding.py
```

#### 5.2.6 Task 6: Test Search

Search for content from your new documents.

### 5.3 Repository

<https://github.com/amornpan/Generic-RAG>

## 5.4 Deliverables

Indexing completed

Documents verified

New documents added

Search tested

**Deadline:** Sunday 23:59

## 6 Lab 6: RAG + Local LLM + Streamlit (3.75%)

### 6.1 Objectives

- Setup Ollama
- Complete RAG pipeline
- Run Streamlit UI

### 6.2 Tasks

#### 6.2.1 Task 1: Install Ollama

1. Download from <https://ollama.ai>
2. Install and run

#### 6.2.2 Task 2: Pull Model

```
ollama pull qwen2.5:7b
```

#### 6.2.3 Task 3: Test LLM

```
ollama run qwen2.5:7b  
# Type: "What is RAG?"
```

#### 6.2.4 Task 4: Run API with LLM

```
cd Generic-RAG  
python api.py
```

#### 6.2.5 Task 5: Run Streamlit

```
streamlit run app.py
```

#### 6.2.6 Task 6: Test Complete Flow

1. Open <http://localhost:8501>
2. Ask questions
3. Verify RAG responses

#### 6.2.7 Task 7: Modify Prompt

Edit the prompt template in `api.py`.

### 6.3 Repository

<https://github.com/amornpan/Generic-RAG>

## 6.4 Deliverables

Ollama running

LLM tested

Complete RAG working

Streamlit UI working

**Deadline:** Sunday 23:59

## 7 Lab 7: Docker Compose (3.75%)

### 7.1 Objectives

- Clone Advanced-RAG-Docker
- Run 6 services with Docker Compose
- Verify all services

### 7.2 Tasks

#### 7.2.1 Task 1: Clone Repository

```
git clone https://github.com/amornpan/Advanced-RAG-Docker.git  
cd Advanced-RAG-Docker
```

#### 7.2.2 Task 2: Study docker-compose.yml

Review the 6 services:

Service	Port
opensearch_container	9200
embedding_container	-
search_api_container	8005
backend_container	8006
frontend_container	8501
ollama_container	11434

#### 7.2.3 Task 3: Start Services

```
docker-compose up -d
```

#### 7.2.4 Task 4: Check Status

```
docker-compose ps
```

#### 7.2.5 Task 5: View Logs

```
docker-compose logs -f
```

#### 7.2.6 Task 6: Test Application

Open <http://localhost:8501>

#### 7.2.7 Task 7: Stop Services

```
docker-compose down
```

### 7.3 Repository

<https://github.com/amornpan/Advanced-RAG-Docker>

## 7.4 Deliverables

All 6 services running

Application accessible

Screenshots of docker ps

Document observations

**Deadline:** Sunday 23:59

## 8 Lab 8: CI/CD + Testing (3.75%)

### 8.1 Objectives

- Write pytest tests
- Create CI/CD pipeline
- Setup GitHub Actions

### 8.2 Tasks

#### 8.2.1 Task 1: Write test\_api.py

```
import pytest
from fastapi.testclient import TestClient
from api import app

client = TestClient(app)

def test_health():
    response = client.get("/health")
    assert response.status_code == 200
    assert response.json()["status"] == "healthy"

def test_search():
    response = client.post("/search",
        json={"query": "test", "top_k": 5})
    assert response.status_code == 200
```

#### 8.2.2 Task 2: Run Tests

```
pytest tests/ -v
pytest tests/ --cov=. --cov-report=html
```

#### 8.2.3 Task 3: Create Jenkinsfile

```
pipeline {
    agent any
    stages {
        stage('Build') {
            steps { sh 'docker-compose build' }
        }
        stage('Test') {
            steps { sh 'pytest tests/' }
        }
        stage('Deploy') {
            steps { sh 'docker-compose up -d' }
        }
    }
}
```

#### 8.2.4 Task 4: Create GitHub Actions

Create `.github/workflows/ci.yml`:

```
name: CI/CD
on: [push, pull_request]
jobs:
  build:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v3
      - run: docker-compose build
      - run: docker-compose up -d
      - run: sleep 60
      - run: curl -f http://localhost:8006/health
```

### 8.3 Repository

<https://github.com/amornpan/Advanced-RAG-Docker>

### 8.4 Deliverables

Tests written

Tests passing

Jenkinsfile created

GitHub Actions workflow

Documentation

**Deadline:** Sunday 23:59

## Summary: All Labs

Lab	Topic	Weight
1	Git + Python Fundamentals	3.75%
2	Docker + OpenSearch	3.75%
3	FastAPI	3.75%
4	OpenSearch Integration	3.75%
5	Embeddings	3.75%
6	RAG + Local LLM + Streamlit	3.75%
7	Docker Compose	3.75%
8	CI/CD + Testing	3.75%
	<b>Total</b>	<b>30%</b>

## Repositories

- Generic-RAG: <https://github.com/amornpan/Generic-RAG>
- Advanced-RAG-Docker: <https://github.com/amornpan/Advanced-RAG-Docker>