

Week 5: OpenSearch Integration

Full Stack RAG with Local LLM

Semester 2/2568

Agenda

1 Vector Search

2 OpenSearch Client

3 Lab 4

Vector vs Keyword Search

Keyword (BM25):

- Exact matching
- Fast
- Miss semantics

Vector Search:

- Semantic similarity
- Understand meaning
- Find related content

Hybrid Search = Best of Both!

Connect to OpenSearch

```
from opensearchpy import OpenSearch

client = OpenSearch(
    hosts=[{"host": "localhost", "port": 9200}] ,
    use_ssl=False
)

# Check connection
info = client.info()
```

Create Index

```
index_body = {
    "settings": {"index": {"knn": True}},
    "mappings": {
        "properties": {
            "content": {"type": "text"},
            "content_vector": {
                "type": "knn_vector",
                "dimension": 1024
            }
        }
    }
}
client.indices.create(index="docs", body=index_body)
```

Lab 4: OpenSearch Integration (3.75%)

Tasks:

- ① Connect from Python
- ② Create index with vector field
- ③ Index sample documents
- ④ Implement hybrid search

Deadline: Sunday 23:59

Questions? See you in Lab 4!